

# **CSE2005L - Operating Systems Lab**

## **Lab 6 - Overhead in system and procedure call**

PRANAVCHENDUR T K - 15BCE1097

Faculty : Dr. Shyamala.L

### **Code to Measure and print System Call**

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/uio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/syscall.h>

// only works on pentium+ x86
// access the pentium cycle counter
// this routine lifted from somewhere on the Web...
void access_counter(unsigned int *hi, unsigned int *lo) {
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1" /* Read cycle counter */
        : "=r" (*hi), "=r" (*lo)           /* and move results to */
        : /* No input */                   /* the two outputs */
        : "%edx", "%eax");
}

// here's the system call we'll use
#define DO_SYSCALL syscall(SYS_getpid)

// calculate difference (in microseconds) between two struct timevals
// assumes difference is less than 2^32 seconds, and unsigned int is 32
bits
unsigned int timediff(struct timeval before, struct timeval after) {
    unsigned int diff;

    diff = after.tv_sec - before.tv_sec;
    diff *= 1000000;
    diff += (after.tv_usec - before.tv_usec);

    return diff;
}

// measure the system call using the cycle counter.  measures the
// difference in time between doing two system calls and doing
// one system call, to try to factor out any measurement overhead
```

```

void measure_cyclecounter(float mhz) {
    unsigned int high_s, low_s, high_e, low_e;
    size_t nbytes;
    float latency_with_read, latency_no_read;

    // warm up all the caches by exercising the functions
    access_counter(&high_s, &low_s);
    // read(5, buf, 4);
    DO_SYSCALL;
    access_counter(&high_e, &low_e);

    // now do it for real
    access_counter(&high_s, &low_s);
    DO_SYSCALL;
    access_counter(&high_e, &low_e);
    latency_with_read = ((float) (low_e - low_s) / mhz);

    access_counter(&high_s, &low_s);
    access_counter(&high_e, &low_e);
    latency_no_read = ((float) (low_e - low_s) / mhz);

    // print out the results
    printf("(cyclecounter) latency: %f microseconds\n", latency_with_read
- latency_no_read);
}

// measure the system call using the cycle counter. measures the
// difference in time between doing two*NLOOPS system calls and doing
// one*NLOOPS system calls, to try to factor out any measurement overhead

#define NUMLOOPS 10000
void measure_gettimeofday() {
    struct timeval beforeone, afterone;
    struct timeval beforetwo, aftertwo;
    int loopcount;
    unsigned int diffone, difftwo, result;

    // warm up all caches
    gettimeofday(&beforeone, NULL);
    gettimeofday(&beforetwo, NULL);
    for (loopcount = 0; loopcount < NUMLOOPS; loopcount++) {
        DO_SYSCALL;
    }
    gettimeofday(&afterone, NULL);

```

```

gettimeofday(&aftertwo, NULL);

// measure loop of one syscall
gettimeofday(&beforeone, NULL);
for (loopcount = 0; loopcount < NUMLOOPS; loopcount++) {
    DO_SYSCALL;
}
gettimeofday(&afterone, NULL);

// measure loop of two syscalls
gettimeofday(&beforetwo, NULL);
for (loopcount = 0; loopcount < NUMLOOPS; loopcount++) {
    DO_SYSCALL;
    DO_SYSCALL;
}
gettimeofday(&aftertwo, NULL);

diffone = timediff(beforeone, afterone);
difftwo = timediff(beforetwo, aftertwo);

result = difftwo - diffone;
printf("(loop) latency:  %f microseconds\n", ((float) result) /
((float) NUMLOOPS));
}

void usage(void) {
    fprintf(stderr, "usage: measure_syscall cpu_mhz\n");
    fprintf(stderr, "  e.g.,  measure_syscall 2791.375\n");
    exit(-1);
}

int main(int argc, char **argv) {

    float mhz;

    if (argc < 2)
        usage();

    if (sscanf(argv[1], "%f", &mhz) != 1)
        usage();

    if ((mhz < 100.0) || (mhz > 100000.0))
        usage();

```

```

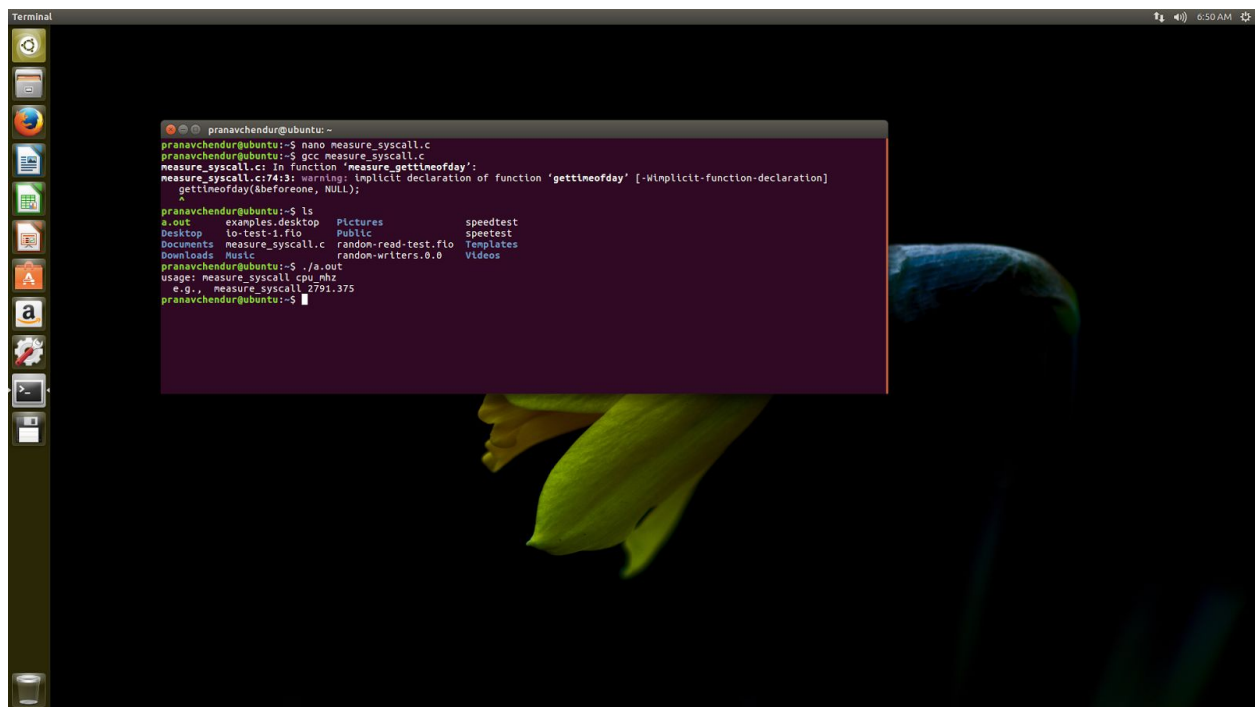
// measure using the cycle counter
measure_cyclecounter(mhz);

// measure using "gettimeofday"
measure_gettimeofday();

return 0;
}

```

## Output



```

pranavchendur@ubuntu:~$ nano measure_syscall.c
pranavchendur@ubuntu:~$ gcc measure_syscall.c
measure_syscall.c: in function 'measure_gettimeofday':
measure_syscall.c:74:3: warning: implicit declaration of function 'gettimeofday' [-Wimplicit-function-declaration]
   gettimeofday(&beforeone, NULL);
   ^
pranavchendur@ubuntu:~$ ls
a.out  examples.desktop  Pictures  speedtest
Desktop  to-test-1.fifo    Public   speetest
Documents  measure_syscall.c  random-read-test.fifo  Templates
Downloads  Music             random-writers.0.0    Videos
pranavchendur@ubuntu:~$ ./a.out
usage: measure_syscall cpu_mhz
e.g., measure_syscall 2791.375
pranavchendur@ubuntu:~$

```

## Code to measure time of System Call Execution

```

#include <sys/time.h>
#include <unistd.h>
#include <assert.h>
#include <stdio.h>

```

```

int foo(){
    return(10);
}

```

```

}

long nanosec(struct timeval t){ /* Calculate nanoseconds in a timeval
structure */
    return((t.tv_sec*1000000+t.tv_usec)*1000);
}

main(){
    int i,j,res;
    long N_iterations=1000000; /* A million iterations */
    float avgTimeSysCall, avgTimeFuncCall;
    struct timeval t1, t2;

    /* Find average time for System call */
    res=gettimeofday(&t1,NULL); assert(res==0);
    for (i=0;i<N_iterations; i++){
        j=getpid();
    }
    res=gettimeofday(&t2,NULL);    assert(res==0);
    avgTimeSysCall = (nanosec(t2) - nanosec(t1))/(N_iterations*1.0);

    /* Find average time for Function call */
    res=gettimeofday(&t1,NULL);    assert(res==0);
    for (i=0;i<N_iterations; i++){
        j=foo();
    }
    res=gettimeofday(&t2,NULL);    assert(res==0);
    avgTimeFuncCall = (nanosec(t2) - nanosec(t1))/(N_iterations*1.0);

    printf("Average time for System call getpid : %f\n",avgTimeSysCall);
    printf("Average time for Function call : %f\n",avgTimeFuncCall);
}

```

**Output:**

```
Terminal
pranavchendur@ubuntu: ~
pranavchendur@ubuntu:~$ nano test_sys_time.c
pranavchendur@ubuntu:~$ gcc test_sys_time.c
test_sys_time.c:15:1: warning: return type defaults to 'int' [-Wimplicit-int]
main() {
^
pranavchendur@ubuntu:~$ ./a.out
Average time for System call getpid : 5.144000
Average time for Function call : 1.970000
pranavchendur@ubuntu:~$
```